Governed AI Analytics Copilot: Roadmap

Draft v1 for review · Team ALTO PALS × Palo Alto Networks Customer Analytics · Prepared Oct 6, 2026 · Demo Day Nov 9, 2026

This roadmap says **when** things happen. **What** we build and why lives in the PRD, which stays stable once Alex approves it. The roadmap is a living document: we update it with each Monday email update to Alex, and timeline changes do not need re-approval of the PRD. Requirement IDs (A1, L2, S1…) refer to the PRD.

1. Milestones at a glance

Milestone	Date	What it means	Who signs off
M0 · Discovery done	Oct 5	Problem, persona and data direction confirmed with Alex	Done
M1 · PRD and roadmap sent	Wed Oct 7	Alex has both documents to review	Approved by Alex Oct 7
M2 · Data and registry	Sun Oct 11	Dataset v1 loaded; 6–10 official metrics answerable	Team
M3 · Consistent and scoped	Sun Oct 18	Any phrasing returns the same number; access matrix passes	Team, shown to Alex Oct 19
M4 · Copilot end to end	Sun Oct 25	Technical seller asks, follows up, sees scope and chart	Shown to Alex Oct 26
M5 · Alerts live	Sun Nov 1	Two triggers deliver privately to the right user	Shown to Alex Nov 2
M6 · Demo ready	Sun Nov 8	All PRD success criteria green; rehearsed twice	Team
Demo Day	Mon Nov 9	Live demo and repo handover	Alex

2. Week by week

Week	Dates	Deliverables (PRD IDs)	Exit check	Depends on
Phase 0	Aug 28–Oct 4	Research, brief analysis, landscape scan, architecture draft, discovery prep	Discovery call held	—
Week 1	Oct 5–Oct 11	PRD and roadmap to Alex (Oct 7) · decide Cube vs. own registry (by Fri Oct 9, after a short spike) · dataset v1: CRM sample plus synthetic engagement and support tables · metric registry, 6–10 metrics · context layer v1: data dictionary, glossary and business rules (per Alex, Oct 7) (A1) · repo, CI and one-command setup	Data loads; registry answers via API; CI runs	Dataset choice; registry tool decision
Week 2	Oct 12–Oct 18	Role scoping in query and database layer (A3) · intent normalization (A2) · phrasing consistency test suite · scope statement (A4) · audit log skeleton (S4)	Same number across 10–20 phrasings; zero out-of-scope rows	M2; mock identities and roles
Week 3	Oct 19–Oct 25	Copilot UI: answer, chart, multi-turn follow-ups with no fixed limit (A5, A6) · injection and social-engineering guardrails v1 (S1, S2) · suspicious-prompt flagging (S3)	Technical-seller flow works end to end	M3
Week 4	Oct 26–Nov 1	Alert engine (L1) · two triggers: high-severity ticket, CXO engagement (L2) · user preferences (L3) · private routing (L4) · Slack or email (L5)	Planted events alert the right person only	M3; synthetic event data; Slack/email test accounts
Week 5	Nov 2–Nov 8	Red-team suite in CI (S6) · repo security scans (S5) · second dataset onboarded (A7) · architecture, deployment and README docs · demo script and two rehearsals	All success criteria green	M4, M5
Demo Day	Nov 9	Live demo · repo handover for Palo Alto Networks infosec review	—	M6

3. Workstreams

Three lanes run in parallel. Each week a rotating DRI keeps the exit check on track. Lane owners are still to be assigned by the team.

Lane	Focus across the 5 weeks	Owner
Data and semantic layer	Dataset, metric registry, intent normalization, role scoping, consistency tests	Mitra, Adedoyin
Copilot and onboarding	Answer flow, follow-ups, charts, UI, second-dataset onboarding	Mofazzel, Trevor
Alerts, governance and security	Alert engine and preferences, private routing, guardrails, audit log, red-team suite, repo scans	Mitra, Adedoyin

4. Client touchpoints

- **Wed Oct 7**: PRD and roadmap sent; Alex approved both (“Looks good”).
  - **Weekly, Monday (email):** a short email update to Alex every Monday, starting Oct 12: what shipped, a screenshot, decisions needed, scope changes and open questions. A call only when needed (per Alex, Oct 7).
  - **Each check-in:** what shipped, a short demo or screenshot, decisions needed, and any scope change with its assumption written down.
  - **Week 5:** walk Alex through tech stack, architecture and deployment, as he asked for near the demo.

5. Dependencies and risks

Risk or dependency	Impact	Plan
Public CRM sample lacks engagement, support or usage data	High	Generate synthetic tables in Week 1 and state it as an assumption
AI still varies across phrasings	High	AI only picks registry metrics; test suite gates every merge; fall back to clarifying questions

Registry tool decision (Cube vs. our own) slips	Medium	Decide by Fri Oct 9 after a short spike; default to our own registry if unclear
Alert scope grows (many roles and triggers)	Medium	Two triggers for Demo Day; others documented as next steps
Slack or email delivery setup blocks testing	Medium	Start with one channel; local mock inbox for tests
No buffer week before Demo Day	High	Week 5 is fixed for hardening; cut scope from Week 4 rather than Week 5

6. Scope-cut order

As Alex suggested, if we run out of time we document the assumption and cut scope instead of shipping half-built features. We cut in this order:

- 1. CSM alert triggers (large deal, low usage)
- 2. Second alert channel (keep one of Slack or email)
- 3. Onboarding UI for new datasets (keep a documented script)
- 4. Charts beyond basic bar and line

Never cut: consistency across phrasings, role scoping below the AI, the audit log and the red-team suite.

7. Change log

Date	Change
Oct 6	v1 drafted from the Oct 5 discovery call
Oct 6	Registry decision moved to Fri Oct 9 so Week 1 is not blocked on delivery day; dates written as month names